

Abstract

- GPGPU-Sim, Multi2Sim, Barra (GPU simulators)
- It can run unmodified OpenCL-based applications

Introduction

- CPU tools
 - Transient faults on hardware structures
 - Noise
 - Hardware reliability
 - OpenRISC
- RTL-level implementations and low-level microarchitecture are lacking
- Goals
 - Realism
 - Flexible
 - Can run OpenCL or CUDA
 - M0, 1, 2 - memory controllers
- Non-goals
 - Not ASIC compatible
- Commonalities between AMD and Nvidia
- Ideas
 - Physical design perspective to “traditional” microarchitecture research
 - Thread-block compaction - warp scheduler technique
 - Circuit-failure prediction
 - Timing speculation
 - Transient fault injection

MIAOW Architecture

- ISA
 - Program counter
 - Execute mask
 - Status registers
 - Mode register
 - General purpose registers

- MIAOW Processor Design Overview
 - Has host CPU
 - Assigns kernel to GPGPU
 - Handled by GPU's ultra threaded dispatcher
 - Computes kernel assignments and schedules wavefronts to CUs
 - Allocates wavefront slots, registers, and LDS space
 - CU execute the kernels
 - scalar and vector ALUs
 - load-store unit
 - **LDS (Local data store) is internal scratch pad memory**
 - have access to device memory through memory controller
 - L1
 - scalar data accesses and instructions
 - Unified L2
 - You can schedule up to **40 wavefronts** on each CU
- Compute Unit Microarchitecture
 - Fetch - interface between dispatcher and compute unit
 - Receives initial PC value
 - Range of registers and local mem it can use
 - Unique ID for wavefront
 - Also informs dispatcher when wavefront is completed
 - Wavepool - Instruction queue for fetched instructions
 - Up to 40 wavefronts - 40 independent queues
 - Can reside in compute unit at a time
 - Decode - instruction
 - Collates 2 32bit-halves of 64bit instructions?
 - *Workgroup - more than 1 wavefront, but mapped to a single CU*
 - Issue/schedule - tracks all in-flight instructions
 - Scoreboard to resolve dependencies
 - Handles barrier and halt instructions
 - Vector ALU - 16 wide (processed in 4 batches)
 - Register files
 - in CU - 1024 vector registers

- 512 scalar registers
 - Vector register files in 64 pages / banks
 - Page corresponds to each thread in wavefront
- Ultra-threaded dispatcher
 - Global wavefront scheduling
 - Receives workgroups from host CPU, passes them to CUs
 - LDS, GDS, and register files never overlap
 - Workgroups come through host interface, if workgroup table available, accept
 - Otherwise, let the host know it can't
 - -> Goes to **resource allocator**
 - finds a CU that has enough resources
 - If free CU found, CU id and allocation data are passed to resource table
 - GPU interface
 - workgroup into wavefronts to the CU
- Design choices
 - Fetch takes in 1 at a time. GCN does 16-32, per cache line